

Laboratory 5: Rigid and Non-Rigid Structure from Motion

Jesus Miguel Adrian Matos

July 2, 2024

Abstract

Nothing for now.

1 Inputs:

The input is a Matrix that one explain below:

$$Input = \begin{pmatrix} I_1 \\ I_2 \\ \vdots \\ I_f \end{pmatrix} \quad (1)$$

Where I_k is a Matrix $2 \times p$, that represent a image in the frame k , such as $p = \text{number of track points}$ and $f = \text{number of frames}$

$$I_k = \begin{pmatrix} x_1 & x_2 & \cdots & x_p \\ y_1 & y_2 & \cdots & y_p \end{pmatrix} \quad (2)$$

2 Outputs:

We want to obtain a representation in the 3D plane of the points tracked in the 2D images.

$$Output = \begin{pmatrix} x_1 & x_2 & \cdots & x_p \\ y_1 & y_2 & \cdots & y_p \\ z_1 & z_2 & \cdots & z_p \end{pmatrix} \quad (3)$$

Where (x, y, z) represents the 3D coordinates in the \mathbb{R}^3 space.

3 First Part (Rigid structure from motion by factorization)

In this task, I must complete the orthogonal factorization function (**rigidfactorization_ortho.m**). Also, we run the code with the provided example (a synthetic sequence with an ogre is provided to you) and report the 3D reconstruction error.

In this task, factorization for a rigid body must be implemented, as shown in the following slide figure 1:

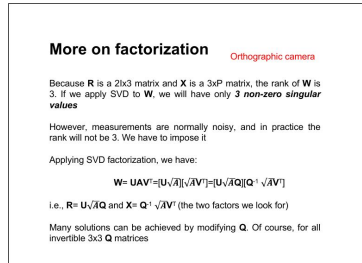


Figure 1: Factoring slide

3.1 Code:

Now we introduce the code, we will explain each line number, but first I want to mention that the value of Zc is not the input matrix, but this is the input matrix minus the average of all the points in each frame. This way we no longer have to consider translation in the calculations.

Therefore the average matrix would be the following (with a dimension of $2f \times p$):

$$Mean = \begin{pmatrix} Mean_1 & Mean_1 & \cdots & Mean_1 \\ Mean_2 & Mean_2 & \cdots & Mean_2 \\ \vdots & \vdots & \ddots & \vdots \\ Mean_f & Mean_f & \cdots & Mean_f \end{pmatrix} \quad (4)$$

where

$$Media_k = \begin{pmatrix} \frac{x_1 + x_2 + \cdots + x_p}{p} \\ \frac{y_1 + y_2 + \cdots + y_p}{p} \end{pmatrix} \quad (5)$$

for any k between 1 and f . Therefore from 4 we have:

$$Zc = Input - Mean \quad (6)$$

Having said this, we continue with the application of the code figure 2.

In the first line, we calculate the factorization called Singular value decomposition(SVD).

In the second line, we approximate $R = U(:, 1 : 3)D(1 : 3, 1 : 3)^{1/2}$ and $S = D(1 : 3, 1 : 3)^{1/2}V(:, 1 : 3)'$ (where R represents rotation and S represents shape) in such a way that we only take the first three columns of U and, the first three columns of V and, the first three rows and the first three columns of D . This is because, being a rigid body, the range of R and S is 3.

In the "while loop" we try to reduce $\|Zc - RU\|$ to the threshold (in the code is epsilon). The idea is to use the rotation condition ($R_i R_i' = I$ where I is identity matrix) which is used to calculate T in the code and thus be able to calculate the new matrices S and R .

```

39         [U, D, V] = svd(Zc);
40         R = U(:, 1:3)*(D(1:3, 1:3) ^ 1/2);
41         S = (D(1:3, 1:3) ^ 1/2)*V(:, 1:3)';
42
43         %-----
44         % Metric Upgrade Step
45         F=size(Zc,1);
46         k=0;
47         thenorm=epsilon+1;
48         while thenorm>epsilon && k<n_iter
49             Rp=R;
50             M=[];
51             for f=1:2:F
52                 Rf=R(f:f+1,:);
53                 [U2,useless,V2]=svd(Rf,'econ');
54                 T=U2*V2';
55                 M=[M;T(1:2,:)];
56             end
57             S=pinv(M)*Zc;
58             R=Zc*pinv(S);
59             k=k+1;
60             thenorm=norm(R-Rp,'fro')/numel(M);
61         end

```

Figure 2: source code

3.2 Examples:

Here we will show how to reconstruct the face of an Ogre with the given frames and using the following input parameters:

- threshold: 10^{-7}
- maximum number of iterations: 200




Figure 3: the original face of the Ogre.




Figure 4: The Ogre's face reconstructed with the algorithm.

As can be seen in Figure 5, the reconstruction of the ogre's face is quite good, the error between the real ogre and the reconstructed one is quite low Figure 6.

4 Second Part (Non-Rigid structure from motion by non-linear optimization and assuming a low-rank shape model)

In this task we must perform a nonlinear optimization for a non-rigid body, assuming a low rank. This code used Levenberg-Marquardt algorithm to optimize the non-linear equation that it is shown to next:

$$\arg \min_{R^i, B_k, l_i^k, t^i} \sum_{i=1}^i \|\widehat{W}^i - R^i \sum_{k=1}^k l_i^k B_k - t^i\|^2 + \gamma \sum_{i=1}^{i-1} \|L^i - L^{i+1}\|^2 + \phi \sum_{i=1}^{i-1} \|R^i - R^{i+1}\|^2 \quad (7)$$

Where: B_k is a row of bases matrix. R^i is a rotacon for a point in column i of X the 3D shape vector. l_i are the constants such that $\sum_{k=1}^k l_i^k B_k = X$. t^i is the translation. L^i is the vector of all constants in frame i . The norm is that of Frobenius.




Figure 5: Comparison between the 3D representation of the Ogre's face against the one reconstructed by the algorithm.

```
Running...
3D error: 0.0026993 %
>>
```

Figure 6: Error between the ogre's face and the reconstructed face.

In order to calculate the Levenberg-Marquardt optimization we need to calculate the Jacobian form of the equation 7, we do this as shown in the following slide Figure 7




Figure 7: Jacobian form

4.1 Code:

Now we will show how this Jacobian form is calculated in the code.

So first we will define the dimensions of Jacobian matrix. For this code we are using quaternions for the rotations, therefore R^i is 2×4 .

Then the dimension of Jacobian matrix is defined in the following way

- number of rows of $Dim(J^{rows}) = 2 \sum_{i=1}^f (numberofpointsvisibleinframei)$

- number of columns of $Dim(J_{columns}) = 4f + 2f + kf + 3kp$

Where f is the number of frames, p is the total number of points tracked and k is the total number of bases.

Explaining row sizes:

Green: Represents the visible points per frame and when we multiply it by two this is because the rotations are in quaternions and the quaternions are 2×4 and then they are two vectors per each points and that is why it is multiplied by two.

Explaining column sizes:

Red: Represents the value of the columns of the rotation matrix (2×4 quaternions) and we have a rotation matrix R^i for each frame i , therefore $4f$ where f is the number of frames.

Blue: Represents the translation matrix t^i , whose dimension is 1×2 , for each frame i , therefore $2f$ where f is the number of frames.

Orange: Represents the matrix L_i whose dimension is $1 \times k$ for each frame i , therefore kf where f is the number of frames.

Yellow: Represents the matrix B_b whose dimension is $3 \times p$ for each base b , therefore $3pk$ where k is the number of bases.

Well then we need to define the shape of our Jacobian matrix by putting 1 where there are values and leaving 0 in the spaces that do not exist, as shown in the following code.

```
for i=1:n_frames
    for j=1:nnz(vij(i,:))
        J(2*j-1+computed_points:2*j+
            computed_points,6*j-5:6*j)=ones(2,6);
        J(2*j-1+computed_points:2*j+
            computed_points,K*j-K+1+6*n_frames:K*j
            +6*n_frames)=ones(2,K);
        J(2*j-1+computed_points:2*j+
            computed_points,3*K*j-3*K+1+(6+K)*
            n_frames:3*K*j+(6+K)*n_frames)=ones
            (2,3*K);
    end
    computed_points=computed_points+2*nnz(vij(i,:))
end
```

In the first for loop each frame is iterated, in the second for loop only the points visible in that frame are iterated.

```
J(2*j-1+computed_points:2*j+computed_points,6*j-5:6*j)
=ones(2,6);
```

Here, the Jacobian is calculated for the rotation R part (in quaternions) and the translation part t , having this

$$(R_{2 \times 4} \mid t_{2 \times 2})_{2 \times 6} \quad (8)$$

we can add a matrix of ones of 2×6

Now, the Jacobian is calculated for the constants L^i part, where $L^i = [l_1^i, l_2^i, \dots, l_k^i]$ such that i is a frame.

```
J(2*j-1+computed_points:2*j+computed_points,K*j-K+1+6*
n_frames:K*j+6*n_frames)=ones(2,K);
```

and for that, in the code is added a matrix of ones from $2 \times k$.

Now, the Jacobian is calculated for the bases B_i part, where

$$B_i = ([b_1^i]_{3 \times 1} \mid [b_2^i]_{3 \times 1} \mid \dots \mid [b_k^i]_{3 \times 1})_{3 \times k} \quad (9)$$

```
J(2*j-1+computed_points:2*j+computed_points,3*K*j-3*K
+1+(6+K)*n_frames:3*K*j+(6+K)*n_frames)=ones(2,3*K)
;
```

and for that, in the code is added a matrix of ones from $2 \times 3k$.

Now we calculate the Jacobian part of the priors, so the dimension of the Jacobian for the priors comes from the following expression

$$\arg \min_{R^i, B_k, l_i^k, t^i} \gamma \sum_{i=1}^{i-1} \|L^i - L^{i+1}\|^2 + \phi \sum_{i=1}^{i-1} \|R^i - R^{i+1}\|^2 \quad (10)$$

so from this equation 10 and the orange color of the equation 7 we deduce that the Jacobian dimension for the prior L would be

$$Dim(J_{rows}) = f - 1Dim(J_{columns}) = kf \quad (11)$$

and the blue and red color of the equation 7 also we can deduce the Jacobian dimension for the prior R as

$$Dim(J_{rows}) = f - 1Dim(J_{columns}) = 6f \quad (12)$$

Now we will show how we fill the indices that have a value in the Jacobian Matrix of the L prior with 1 and otherwise hold with zero.

```

if (priors.coeff_prior == 1)
    % prior terms on L
    L=spalloc(n_frames-1,(K+6)*n_frames + K*3*
        n_points,2*K*(n_frames-1));
    for i=1:n_frames-1
        L(i,K*i-K+1+6*n_frames:K*i+K+6*n_frames)=ones
            (1,2*K);
    end
    if K==1
        J(2*nnz(vij)+1:size(J,1), :) = L;
    else
        J(2*nnz(vij)+1:2*nnz(vij)+n_frames-1, :) = L;
    end
end
end

```

We do the following, we take the L_i and L_{i+1} and join them as shown in the equation 13

$$\left([L_i]_{1 \times k} \mid [L_{i+1}]_{1 \times k} \right)_{1 \times 2k} \quad (13)$$

and because of this in the code we add a matrix of ones $1 \times 2K$.

For the Jacobian Matrix of the R prior is the similar way, in this case we joint R^i and R^{i+1} like show in the equation 14

$$\left([R^i]_{1 \times 6} \mid [R_{i+1}]_{1 \times 6} \right)_{1 \times 12} \quad (14)$$

and because of this in the code we add a matrix of ones 1×12 .

Joining the Jacobian and the Jacobian of the priors we have the following graph figure 8.




Figure 8: Jacobian result

4.2 Examples:

Here we will show how to reconstruct the human face with the given frames and using the following input parameters:

- low-rank K (it is the number of bases): 2
- $camera_{prior}$ (it mean that the R prior is active): 1
- $coeff_{prior}$ (it mean that the L prior is active): 1
- number of optimizer iterations: 50

In Figure 9 we can see the comparison of the input points of a frame in green against the estimated points of that frame in blue.




Figure 9: Input points compared to optimized points

In Figure 10 we can see the estimated 3D points for a frame.

In Figure 11 we see the calculation of each of the 50 optimizations, this method has a high cost in performance, but a good result is achieved as shown by the low error at the end.




Figure 10: 3D reconstruction

Iteration	Func-count	Resnorm	First-order optimality	Lambda	Norm of step
0	1225	1.84649e+06	9.51e+04	0.01	
1	2450	26864.9	2.09e+03	0.001	25.0136
2	3678	16728	1.07e+03	1	15.9044
3	4903	12361	487	0.1	24.5625
4	6129	11207.5	837	1	15.9301
5	7354	10545.2	661	0.1	12.7065
6	8580	10245.9	717	1	7.64863
7	9805	10085.3	465	0.1	5.64436
8	11031	9981.24	272	1	4.18336
9	12256	9910.45	164	0.1	3.76857
10	13482	9851.74	138	1	3.28008
11	14707	9806.24	108	0.1	2.81746
12	15932	9769.25	165	0.01	5.31267
13	17158	9728.06	114	0.1	3.42729
14	18383	9704.76	96.1	0.01	3.27792
15	19609	9687.54	81.7	0.1	2.7978
16	20834	9675.29	48.5	0.01	2.4206
17	22060	9666.78	42.2	0.1	2.06087
18	23285	9660.79	36.6	0.01	1.83145
19	24511	9656.37	30.5	0.1	1.78
20	25736	9652.86	25.5	0.01	1.81386
21	26962	9649.82	22.3	0.1	1.9704
22	28187	9646.98	20.2	0.01	2.11821
23	29413	9644.15	19.7	0.1	2.32606
24	30638	9641.23	19.1	0.01	2.48918
25	31864	9638.15	19.8	0.1	2.69185
26	33089	9634.88	19.7	0.01	2.83915
27	34315	9631.4	21	0.1	3.02158
28	35540	9627.71	20.9	0.01	3.14215
29	36766	9623.8	22.1	0.1	3.29737
30	37991	9619.73	22.9	0.01	3.38541
31	39217	9615.48	22.4	0.1	3.50923
32	40442	9611.11	24.7	0.01	3.56288
33	41668	9606.61	22.1	0.1	3.65309
34	42893	9602.06	26.4	0.01	3.67613
35	44119	9597.42	21.1	0.1	3.73425
36	45344	9592.76	28	0.01	3.73824
37	46570	9588.01	20.6	0.1	3.76665
38	47795	9583.3	29.4	0.01	3.76896
39	49021	9578.49	23.2	0.1	3.76972
40	50246	9573.86	40.4	0.01	3.79026
41	51472	9569.16	30.5	0.1	3.76321
42	52697	9565.07	73.9	0.01	3.81738
43	53923	9560.89	73.9	0.1	3.76668
44	55148	9558.3	153	0.01	3.86225
45	56374	9555.13	168	0.1	3.80734
46	57601	9550.17	104	10	0.546409
47	58826	9547.64	80.2	1	0.302035
48	60051	9546.47	34.9	0.1	0.763707
49	61276	9542.47	44.1	0.01	3.79643
50	62502	9539.54	86.7	0.1	3.62819

[Solver stopped prematurely.](#)

lsqnonlin stopped because it exceeded the iteration limit,
[options.MaxIterations](#) = 5.000000e+01.

Final reprojection (error: 0.0030556%)

Figure 11: Optimizer iterations